

- **Benchmark Usage Guide**

Measure the performance of your Fake User Generator.

- **Quick Start**

Run Full Benchmark Suite

```
python benchmark.py full
```

This runs all tests and saves results to benchmark_results.json.

Output Example:

```
=====
```

FAKE USER GENERATOR BENCHMARK

```
=====
```

Started at: 2024-12-10 15:30:45

```
=====
```

==== TEST 1: QUICK PERFORMANCE TEST ===

```
=====
```

Benchmarking: 100 users x 5 iterations

Locale: en_US, Seed: 12345

```
=====
```

Iteration 1: 45.23ms | 2,211 users/sec | 100 users generated

Iteration 2: 43.87ms | 2,279 users/sec | 100 users generated

Iteration 3: 44.12ms | 2,266 users/sec | 100 users generated

Iteration 4: 43.95ms | 2,275 users/sec | 100 users generated

Iteration 5: 44.31ms | 2,257 users/sec | 100 users generated

SUMMARY STATISTICS

Average throughput: 2,258 users/sec

Peak throughput: 2,279 users/sec

...

- **Individual Tests**

Quick Performance Test

`python benchmark.py quick`

Tests generation of 100 users over 10 iterations. Good for quick validation.

Scaling Test

`python benchmark.py scaling`

Tests batch sizes: 1, 5, 10, 25, 50, 100, 250, 500, 1000 users.

Output:

SCALING COMPARISON

Batch Size | Avg Time | Throughput | Time/User

| Batch Size | Avg Time | Throughput | Time/User |
|------------|----------|------------|-----------|
| 1 | 5.23ms | 191 u/s | 5.230ms |
| 5 | 12.45ms | 402 u/s | 2.490ms |

| | | | |
|------|------------|---------|---------|
| 10 | 18.67ms | 536 u/s | 1.867ms |
| 25 | 35.12ms | 712 u/s | 1.405ms |
| 50 | 62.34ms | 802 u/s | 1.247ms |
| 100 | 118.45ms | 844 u/s | 1.185ms |
| 250 | 287.23ms | 870 u/s | 1.149ms |
| 500 | 565.12ms | 885 u/s | 1.130ms |
| 1000 | 1,124.56ms | 889 u/s | 1.125ms |

Findings: Time per user decreases as batch size increases (amortized overhead).

Locale Comparison

python benchmark.py locales

Compares performance across different locales.

Output:

LOCALE COMPARISON

| Locale | Avg Time | Throughput |
|--------|----------|------------|
| en_US | 118.2ms | 846 u/s |
| de_DE | 121.5ms | 823 u/s |

Individual Function Benchmark

python benchmark.py functions

Measures performance of each stored procedure individually.

Output:

INDIVIDUAL FUNCTION BENCHMARK

| | |
|------------------------------|--|
| prng_int | Avg: 0.045ms Med: 0.043ms 22,222 calls/sec |
| prng_float | Avg: 0.047ms Med: 0.045ms 21,277 calls/sec |
| prng_normal | Avg: 0.089ms Med: 0.086ms 11,236 calls/sec |
| prng_sphere_coords | Avg: 0.092ms Med: 0.089ms 10,870 calls/sec |
| generate_name | Avg: 2.345ms Med: 2.312ms 426 calls/sec |
| generate_address | Avg: 1.876ms Med: 1.845ms 533 calls/sec |
| generate_phone | Avg: 0.234ms Med: 0.228ms 4,274 calls/sec |
| generate_email | Avg: 0.567ms Med: 0.551ms 1,764 calls/sec |
| generate_physical_attributes | Avg: 1.123ms Med: 1.098ms 891 calls/sec |

- **Understanding Results**

Key Metrics

Throughput (users/sec):

- How many users can be generated per second
- Higher is better
- Example: 2,000 users/sec means 1M users in ~8 minutes

Latency (ms):

- Time to generate one batch
- Lower is better
- Example: 50ms for 100 users = 0.5ms per user

Time/User:

- Average time to generate one user
- Decreases with larger batches (overhead amortization)

- Example: 1.2ms per user at batch size 100

Performance Expectations

Local Development:

- 1,000 - 3,000 users/sec
- 50-150ms for 100 users

Cloud Deployment (shared CPU):

- 500 - 1,500 users/sec
- 100-250ms for 100 users

Cloud Deployment (dedicated CPU):

- 2,000 - 5,000 users/sec
 - 30-80ms for 100 users
-

- **Optimization Tips**

If Performance is Slow

1. Check Database Connection:

-- Check for slow queries

```
SELECT query, calls, mean_exec_time, max_exec_time
FROM pg_stat_statements
WHERE query LIKE '%generate_fake_users%'
ORDER BY mean_exec_time DESC;
```

2. Database Tuning:

-- Increase work memory for complex queries

```
SET work_mem = '64MB';
```

```
-- Check current settings
```

```
SHOW shared_buffers;
```

```
SHOW work_mem;
```

3. Index Optimization:

```
-- Ensure indexes exist
```

```
\d+ first_names
```

```
\d+ last_names
```

```
-- Add missing indexes
```

```
CREATE INDEX IF NOT EXISTS idx_first_names_locale_gender
```

```
ON first_names(locale_code, gender);
```

4. Connection Pooling:

```
# Use connection pooling in app.py
```

```
from psycopg2 import pool
```

```
connection_pool = psycopg2.pool.SimpleConnectionPool(
```

```
    1, 20,
```

```
    **DB_CONFIG
```

```
)
```

- **Comparing Hardware**

Run benchmark on different systems:

Local Machine:

```
python benchmark.py quick > benchmark_local.txt
```

Cloud Server:

```
python benchmark.py quick > benchmark_cloud.txt
```

Compare:

Windows

```
Compare-Object (Get-Content benchmark_local.txt) (Get-Content benchmark_cloud.txt)
```

- **Sample Benchmark Results**

High-Performance Setup

(i7-10700K, 32GB RAM, NVMe SSD, PostgreSQL 14 local)

Batch Size: 100 users

Iterations: 10

Average time: 42.3ms

Throughput: 2,364 users/sec

Time per user: 0.42ms

Cloud Setup

(Render \$7/month, PostgreSQL shared)

Batch Size: 100 users

Iterations: 10

Average time: 187.5ms

Throughput: 533 users/sec

Time per user: 1.88ms

Raspberry Pi 4

(4GB RAM, PostgreSQL 13)

Batch Size: 100 users

Iterations: 10

Average time: 342.8ms

Throughput: 292 users/sec

Time per user: 3.43ms

- **Continuous Benchmarking**

Track Performance Over Time

```
# Run benchmark daily
```

```
$timestamp = Get-Date -Format "yyyy-MM-dd"
```

```
python benchmark.py full
```

```
Rename-Item benchmark_results.json "benchmark_$timestamp.json"
```

Compare Results

```
import json
```

```
import glob
```

```
files = sorted(glob.glob('benchmark_*.json'))
```

```
print("Date      | Throughput | Avg Latency")
```

```
print("-----+-----+-----")
```

```
for file in files:
```

```
    with open(file) as f:
```

```
        data = json.load(f)
```

```
        date = file.replace('benchmark_', '').replace('.json', '')
```

```
        throughput = data['quick_test']['avg_users_per_sec']
```

```
        latency = data['quick_test']['avg_time_ms']
```

```
        print(f'{date} | {throughput:>6.0f} u/s | {latency:>7.2f}ms")
```

- **Benchmarking Different Configurations**

Test with Different Batch Sizes

```
# Custom benchmark
```

```
from benchmark import benchmark_generation
```

```
for size in [10, 50, 100, 500, 1000]:
```

```
    print(f"\n{'='*50}")
```

```
    print(f"Testing batch size: {size}")
```

```
    benchmark_generation('en_US', 12345, size, 5)
```

Test with Different Locales

```
for locale in ['en_US', 'de_DE']:
```

```
    print(f"\n{'='*50}")
```

```
    print(f"Testing locale: {locale}")
```

```
    benchmark_generation(locale, 12345, 100, 10)
```

- **Interpreting Bottlenecks**

Database is Bottleneck

Symptoms:

- High latency even for small batches
- CPU usage low
- Network usage low

Solutions:

- Upgrade database tier
- Enable connection pooling
- Optimize queries

Network is Bottleneck

Symptoms:

- Large difference between local and cloud benchmarks
- High latency variance

Solutions:

- Deploy app closer to database (same region)
- Use internal networking (not public internet)
- Enable compression

Application is Bottleneck

Symptoms:

- High CPU usage
- Latency increases linearly with batch size

Solutions:

- Scale horizontally (more instances)
 - Upgrade CPU
 - Enable caching
-

- **Production Monitoring**

Add to Your App

```
# In app.py
```

```
import time
```

```
@app.route('/api/generate', methods=['POST'])
```

```
def api_generate():
```

```
    start = time.time()
```

```
# ... existing code ...
```

```
    elapsed = time.time() - start
```

```
    app.logger.info(f"Generated {batch_size} users in {elapsed*1000:.2f}ms")
```

```
    return jsonify({
```

```
        'users': users,
```

```
        'generation_time_ms': elapsed * 1000,
```

```
        # ... other fields
```

```
    })
```

Dashboard Monitoring

Most platforms provide:

- Request duration
- Requests per second
- Database query time
- Error rate

Set up alerts for:

- Latency > 500ms
- Error rate > 1%
- Database CPU > 80%

- **Expected Results Summary**

| Batch Size | Time (Local) | Time (Cloud) | Throughput |
|-------------------|---------------------|---------------------|-------------------|
| 10 | 15-25ms | 50-100ms | 400-600 u/s |
| 100 | 40-80ms | 150-300ms | 800-2000 u/s |
| 1000 | 400-800ms | 1-3s | 1000-2500 u/s |

Conclusion: Batch size 100-250 offers best balance of latency and throughput for most use cases.